

# TRANSACTIONS AND CONCURRENCY CONTROL



Edit with WPS Office

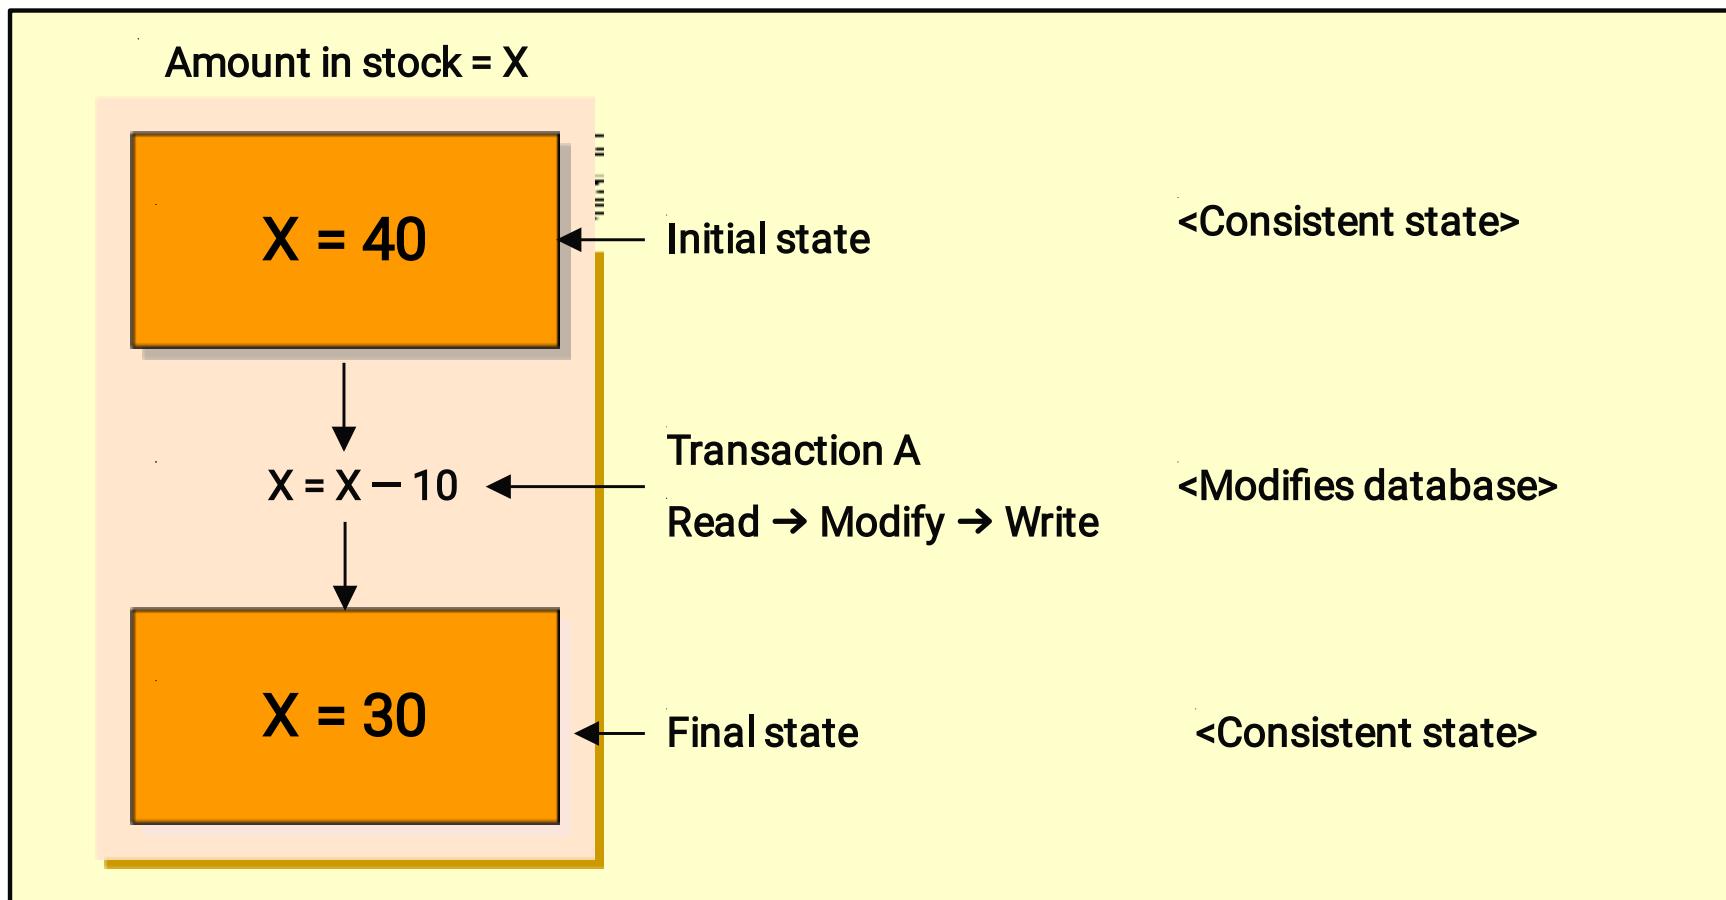
# WHAT IS A TRANSACTION?

- ☒ Logical unit of work that must be either entirely completed or aborted
- ☒ Successful transaction changes database from one consistent state to another
  - ☒ One in which all data integrity constraints are satisfied
- ☒ Most real-world database transactions are formed by two or more database requests
  - ☒ Equivalent of a single SQL statement in an application program or transaction



Edit with WPS Office

# TRANSACTION - EXAMPLE



Edit with WPS Office

# EVALUATING TRANSACTION RESULTS

- ☒ Not all transactions update database
- ☒ SQL code represents a transaction because database was accessed
- ☒ Improper or incomplete transactions can have devastating effect on database integrity
  - ☒ User can define enforceable constraints based on business rules
  - ☒ Entity and referential integrity rules are enforced automatically by the DBMS



Edit with WPS Office

# TRANSACTION PROPERTIES

## ☒ Atomicity

- ☒ All operations of a transaction must be completed

## ☒ Consistency

- ☒ Permanence of database's consistent state

## ☒ Isolation

- ☒ Data used during transaction cannot be used by second transaction until the first is completed



Edit with WPS Office

# TRANSACTION PROPERTIES (CONT'D.)

- ☒ Durability
  - ☒ Once transactions are committed, they cannot be undone
- ☒ Serializability
  - ☒ Concurrent execution of several transactions yields consistent results
- ☒ Multiuser databases are subject to multiple concurrent transactions



Edit with WPS Office

# TRANSACTION MANAGEMENT WITH SQL

- ☒ ANSI has defined standards that govern SQL database transactions
- ☒ Transaction support is provided by two SQL statements: COMMIT and ROLLBACK
- ☒ Transaction sequence must continue until:
  - ☒ COMMIT statement is reached
  - ☒ ROLLBACK statement is reached
  - ☒ End of program is reached
  - ☒ Program is abnormally terminated



Edit with WPS Office

# THE TRANSACTION LOG

- ☒ Transaction log stores:
  - ☒ A record for the beginning of transaction
  - ☒ For each transaction component:
    - ☐ Type of operation being performed (update, delete, insert)
    - ☐ Names of objects affected by transaction
    - ☐ “Before” and “after” values for updated fields
    - ☐ Pointers to previous and next transaction log entries for the same transaction
  - ☒ Ending (COMMIT) of the transaction



Edit with WPS Office

# TRANSACTION LOG - EXAMPLE

TABLE  
13.1

A Transaction Log

TRL_ID	TRX_NUM	PREV_PTR	NEXT_PTR	OPERATION	TABLE	ROW ID	ATTRIBUTE	BEFORE VALUE	AFTER VALUE
341	101	Null	352	START	****Start Transaction				
352	101	341	363	UPDATE	PRODUCT	1558-QW1	PROD_QOH	25	23
363	101	352	365	UPDATE	CUSTOMER	10011	CUST_BALANCE	525.75	615.73
365	101	363	Null	COMMIT	**** End of Transaction				



TRL\_ID = Transaction log record ID

TRX\_NUM = Transaction number

PTR = Pointer to a transaction log record ID

(Note: The transaction number is automatically assigned by the DBMS.)

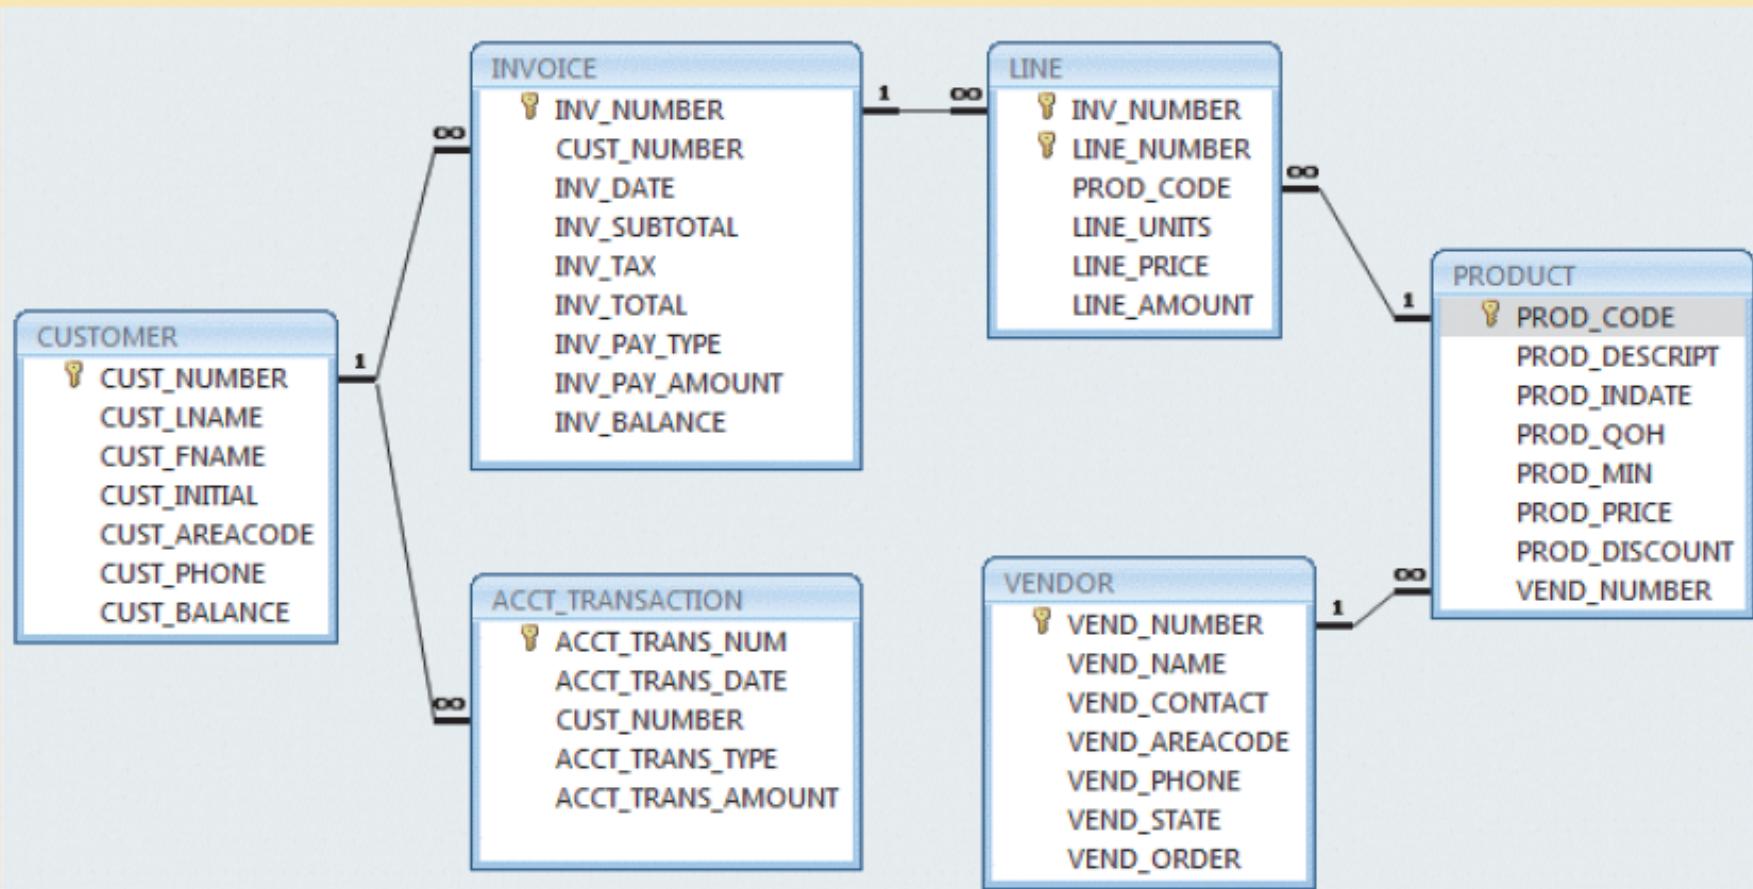


Edit with WPS Office

# SAMPLE DATABASE FOR DISCUSSION

FIGURE  
13.1

The Ch13\_SaleCo database relational diagram



Edit with WPS Office

SOURCE: Course Technology/Cengage Learning

# CONCURRENCY CONTROL

- ☒ Coordination of simultaneous transaction execution in a multiprocessing database
- ☒ Objective is to ensure serializability of transactions in a multiuser environment
- ☒ Three main problems:
  - ☒ Lost updates
  - ☒ Uncommitted data
  - ☒ Inconsistent retrievals



Edit with WPS Office

# LOST UPDATES

- ☒ Lost update problem:
  - ☒ Two concurrent transactions update same data element
  - ☒ One of the updates is lost
    - ☐ Overwritten by the other transaction



Edit with WPS Office

# LOST UPDATE - EXAMPLE

TABLE  
13.2

## Two Concurrent Transactions to Update QOH

TRANSACTION	COMPUTATION
T1: Purchase 100 units	$\text{PROD\_QOH} = \text{PROD\_QOH} + 100$
T2: Sell 30 units	$\text{PROD\_QOH} = \text{PROD\_QOH} - 30$

TABLE  
13.3

## Serial Execution of Two Transactions

TIME	TRANSACTION	STEP	STORED VALUE
1	T1	Read PROD_QOH	35
2	T1	$\text{PROD\_QOH} = 35 + 100$	
3	T1	Write PROD_QOH	135
4	T2	Read PROD_QOH	135
5	T2	$\text{PROD\_QOH} = 135 - 30$	
6	T2	Write PROD_QOH	105



Edit with WPS Office

# LOST UPDATE - EXAMPLE (CONT')

TABLE  
13.4

Lost Updates

TIME	TRANSACTION	STEP	STORED VALUE
1	T1	Read PROD_QOH	35
2	T2	Read PROD_QOH	35
3	T1	$\text{PROD\_QOH} = 35 + 100$	
4	T2	$\text{PROD\_QOH} = 35 - 30$	
5	T1	Write PROD_QOH (Lost update)	135
6	T2	Write PROD_QOH	5

Another way of illustrating lost update for concurrent transaction T1 and T2:

Time	T1	T2	PROD_QOH
t1	begin_transaction		35
t2	read(PROD_QOH)	begin_transaction	35
t3	$\text{PROD\_QOH} = \text{PROD\_QOH} + 100$	read(PROD_QOH)	35
t4	write(PROD_QOH)	$\text{PROD\_QOH} = \text{PROD\_QOH} - 30$	135
t5	COMMIT	write(PROD_QOH)	5
t6		COMMIT	5



Edit with WPS Office

# UNCOMMITTED DATA

- ☒ Uncommitted data phenomenon:
  - ☒ Two transactions are executed concurrently
  - ☒ First transaction rolled back after second already accessed uncommitted data



Edit with WPS Office

# UNCOMMITTED DATA - EXAMPLE

TABLE  
13.5

Transactions Creating an Uncommitted Data Problem

TRANSACTION	COMPUTATION
T1: Purchase 100 units	$\text{PROD\_QOH} = \text{PROD\_QOH} + 100$ (Rolled back)
T2: Sell 30 units	$\text{PROD\_QOH} = \text{PROD\_QOH} - 30$

Time	T1	T2	PROD_QOH
t1	begin_transaction		35
t2	read(PROD_QOH)		35
t3	$\text{PROD\_QOH} = \text{PROD\_QOH} + 100$		35
t4	write(PROD_QOH)		135
t5	ROLLBACK	begin_transaction	35
t6		read(PROD_QOH)	35
t7		$\text{PROD\_QOH} = \text{PROD\_QOH} - 30$	35
t8		write(PROD_QOH)	5
T9		COMMIT	5



Edit with WPS Office

# UNCOMMITTED DATA - EXAMPLE (CONT')

## ☒ Uncommitted data

Time	T1	T2	PROD_QOH
t1	begin_transaction		35
t2	read(PROD_QOH)		35
t3	PROD_QOH = PROD_QOH +100	begin_transaction	35
t4	write(PROD_QOH)	read(PROD_QOH)	135
t5	ROLLBACK	PROD_QOH = PROD_QOH - 30	135
t6		write(PROD_QOH)	105
t7		COMMIT	105



Edit with WPS Office

# INCONSISTENT RETRIEVALS

- ☒ Inconsistent retrievals:
  - ☒ First transaction accesses data
  - ☒ Second transaction alters the data
  - ☒ First transaction accesses the data again
- ☒ Transaction might read some data before they are changed and other data after changed
- ☒ Yields inconsistent results



Edit with WPS Office

# INCONSISTENT RETRIEVAL -

## EXAMPLE

TABLE  
13.8

### Retrieval During Update

TRANSACTION T1	TRANSACTION T2
SELECT SUM(PROD_QOH) FROM PRODUCT	UPDATE PRODUCT SET PROD_QOH = PROD_QOH + 10 WHERE PROD_CODE = 1546-QQ2
	UPDATE PRODUCT SET PROD_QOH = PROD_QOH - 10 WHERE PROD_CODE = 1558-QW1
	COMMIT;

TABLE  
13.9

### Transaction Results: Data Entry Correction

	BEFORE	AFTER
PROD_CODE	PROD_QOH	PROD_QOH
11QER/31	8	8
13-Q2/P2	32	32
1546-QQ2	15	(15 + 10) → 25
1558-QW1	23	(23 - 10) → 13
2232-QTY	8	8
2232-QWE	6	6
Total	92	92



Edit with WPS Office

# INCONSISTENT RETRIEVAL - EXAMPLE (CONT')

TABLE  
13.10

Inconsistent Retrievals

TIME	TRANSACTION	ACTION	VALUE	TOTAL
1	T1	Read PROD_QOH for PROD_CODE = '11QER/31'	8	8
2	T1	Read PROD_QOH for PROD_CODE = '13-Q2/P2'	32	40
3	T2	Read PROD_QOH for PROD_CODE = '1546-QQ2'	15	
4	T2	PROD_QOH = 15 + 10		
5	T2	Write PROD_QOH for PROD_CODE = '1546-QQ2'	25	
6	T1	Read PROD_QOH for PROD_CODE = '1546-QQ2'	25	(After) 65
7	T1	Read PROD_QOH for PROD_CODE = '1558-QW1'	23	(Before) 88
8	T2	Read PROD_QOH for PROD_CODE = '1558-QW1'	23	
9	T2	PROD_QOH = 23 – 10		
10	T2	Write PROD_QOH for PROD_CODE = '1558-QW1'	13	
11	T2	***** COMMIT *****		
12	T1	Read PROD_QOH for PROD_CODE = '2232-QTY'	8	96
13	T1	Read PROD_QOH for PROD_CODE = '2232-QWE'	6	102



Edit with WPS Office

# THE SCHEDULER

- ✉ Special DBMS program
  - ✉ Purpose is to establish order of operations within which concurrent transactions are executed
- ✉ Interleaves execution of database operations:
  - ✉ Ensures serializability
  - ✉ Ensures isolation
- ✉ Serializable schedule
  - ✉ Interleaved execution of transactions yields same results as serial execution
- ✉ Two concurrent transactions are in conflict when
  - ✉ they access the same data, and
  - ✉ at least one of them executes WRITE operation



Edit with WPS Office

# CONCURRENCY CONTROL WITH LOCKING METHODS

## ☒ Lock

- ☒ Guarantees exclusive use of a data item to a current transaction
- ☒ Required to prevent another transaction from reading inconsistent data
- ☒ Pessimistic locking
  - Use of locks based on the assumption that conflict between transactions is likely

## ☒ Lock manager

- Responsible for assigning and policing the locks used by transactions



Edit with WPS Office

# LOCK GRANULARITY

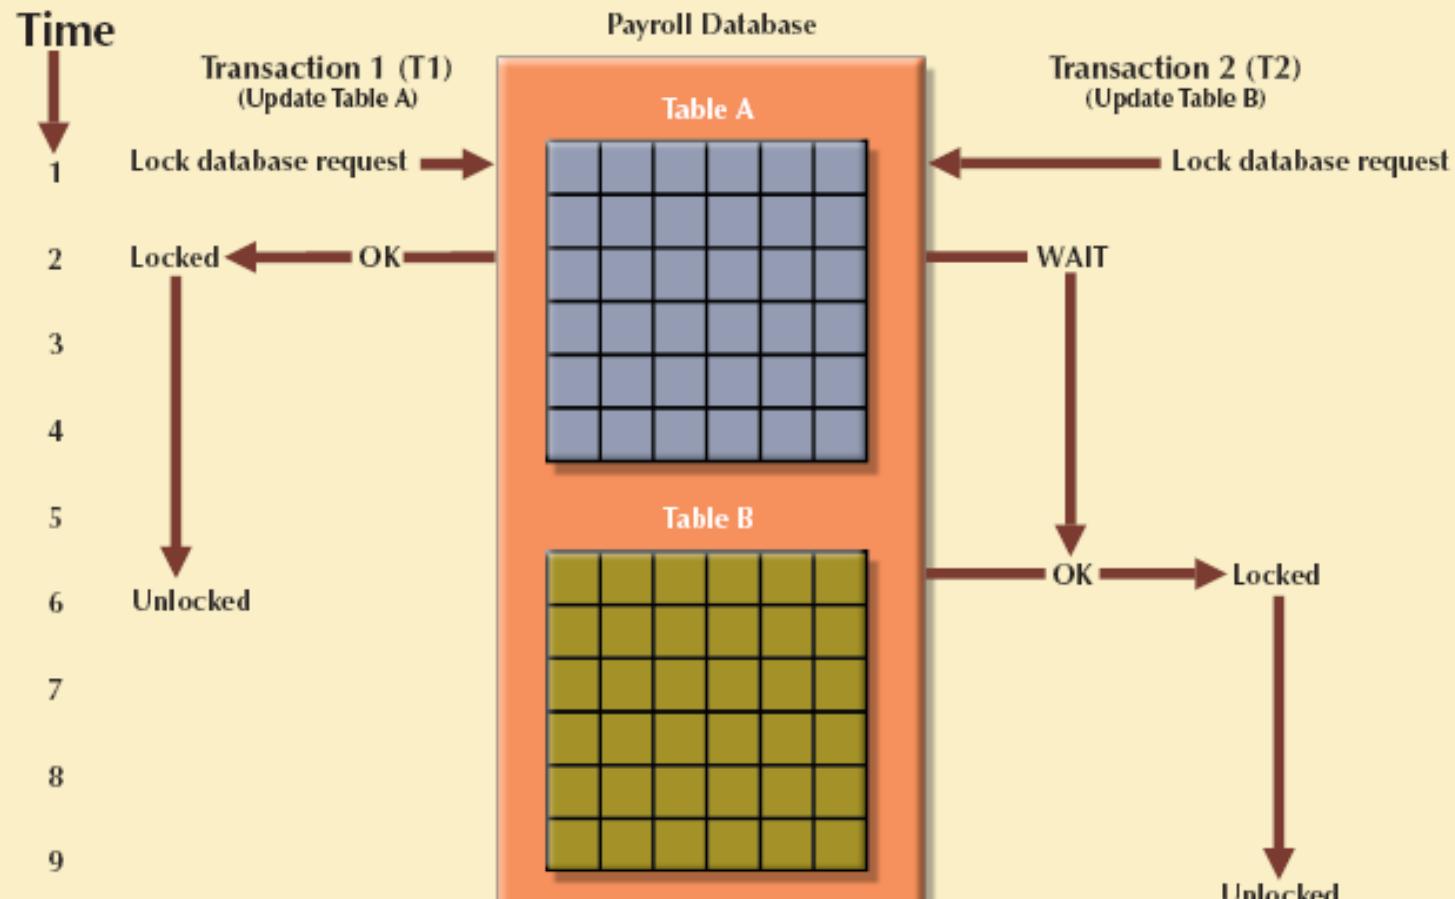
- ☒ Indicates level of lock use
- ☒ Locking can take place at following levels:
  - ☒ Database - Entire database is locked
  - ☒ Table - Entire table is locked
  - ☒ Page - Entire diskpage is locked
  - ☒ Row-level lock
    - ☐ Allows concurrent transactions to access different rows of same table
    - ☐ Even if rows are located on same page
  - ☒ Field-level lock
    - ☐ Allows concurrent transactions to access same row as long as they are different fields (attributes) within the row



Edit with WPS Office

# DATABASE-LEVEL LOCK

FIGURE  
13.3 Database-level locking sequence

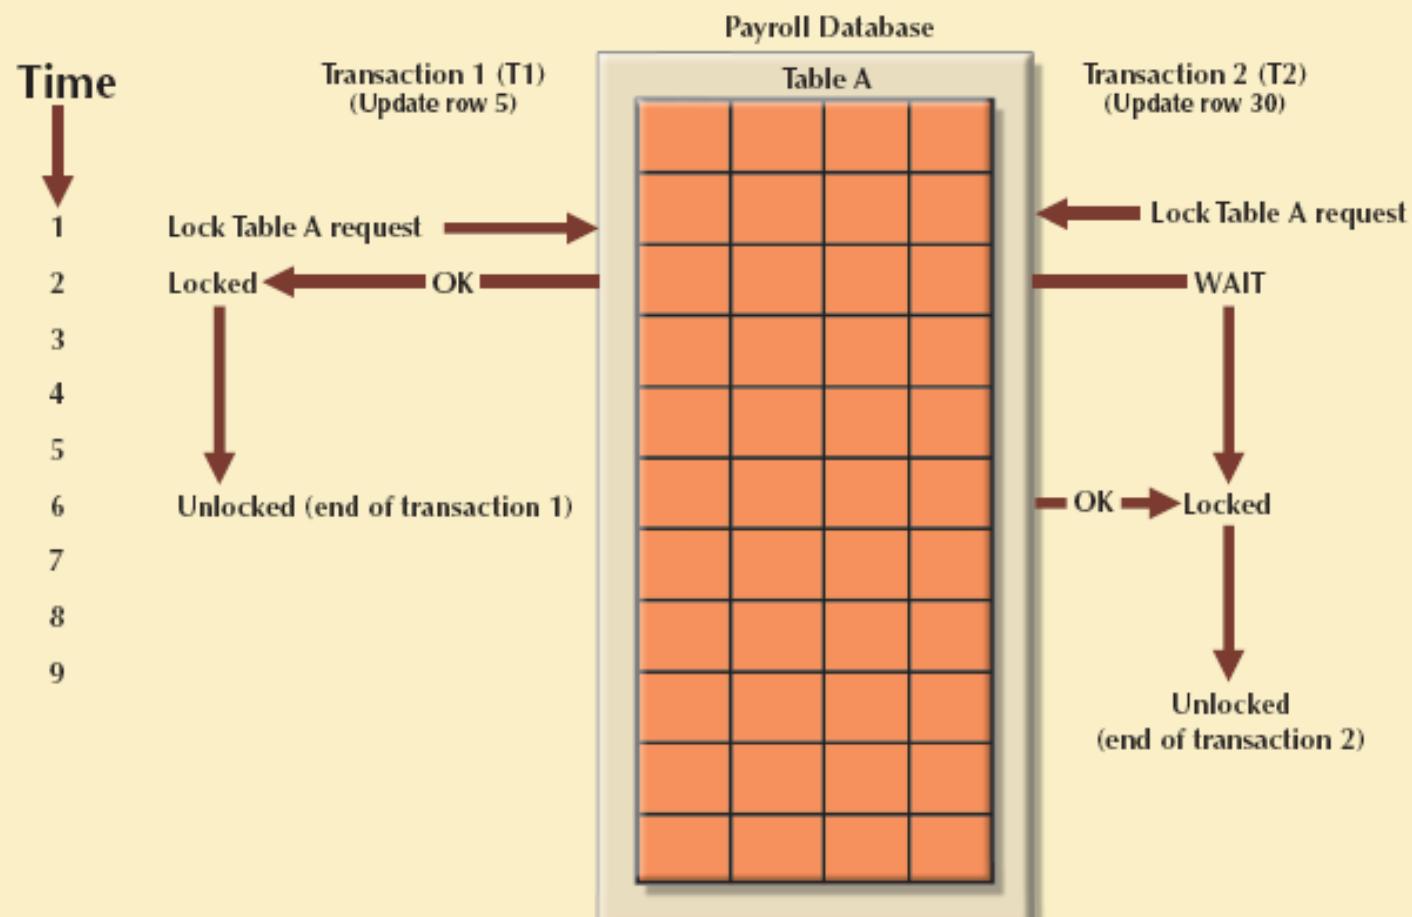


Edit with WPS Office

SOURCE: Course Technology/Cengage Learning

# TABLE-LEVEL LOCK

FIGURE 13.4 An example of a table-level lock



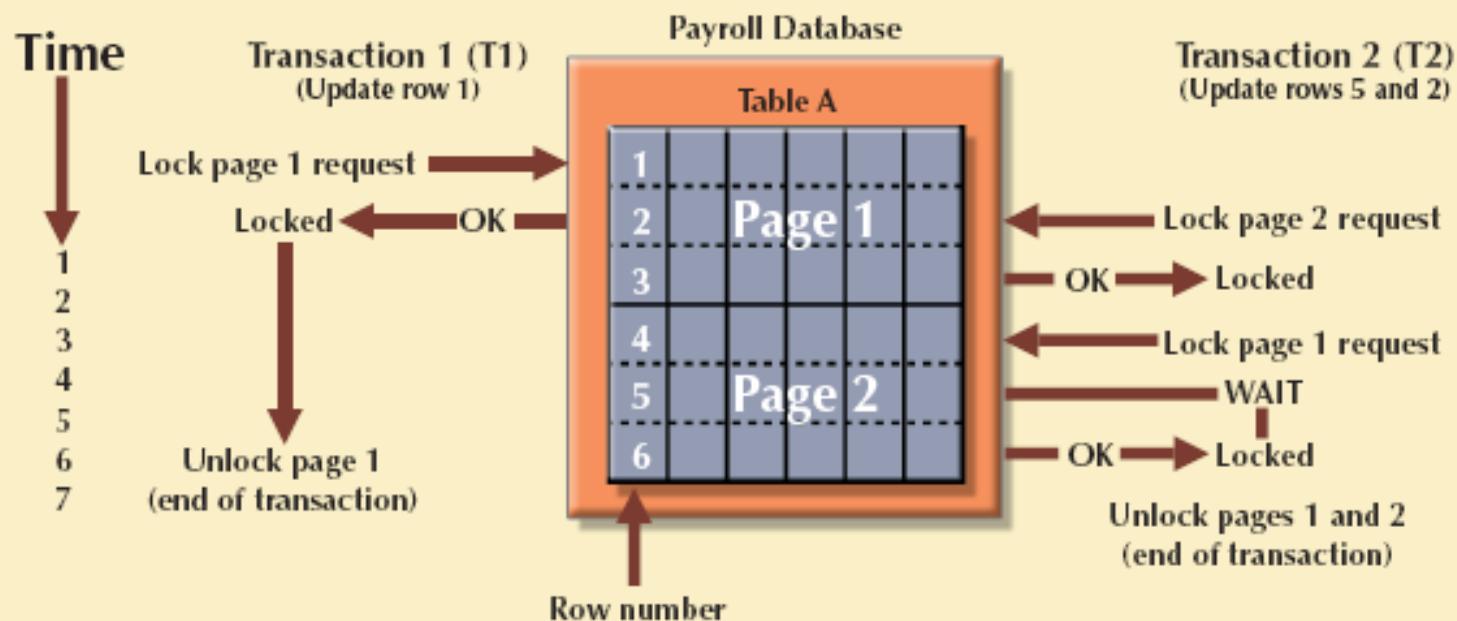
SOURCE: Course Technology/Cengage Learning



Edit with WPS Office

# PAGE-LEVEL LOCK

FIGURE 13.5 An example of a page-level lock



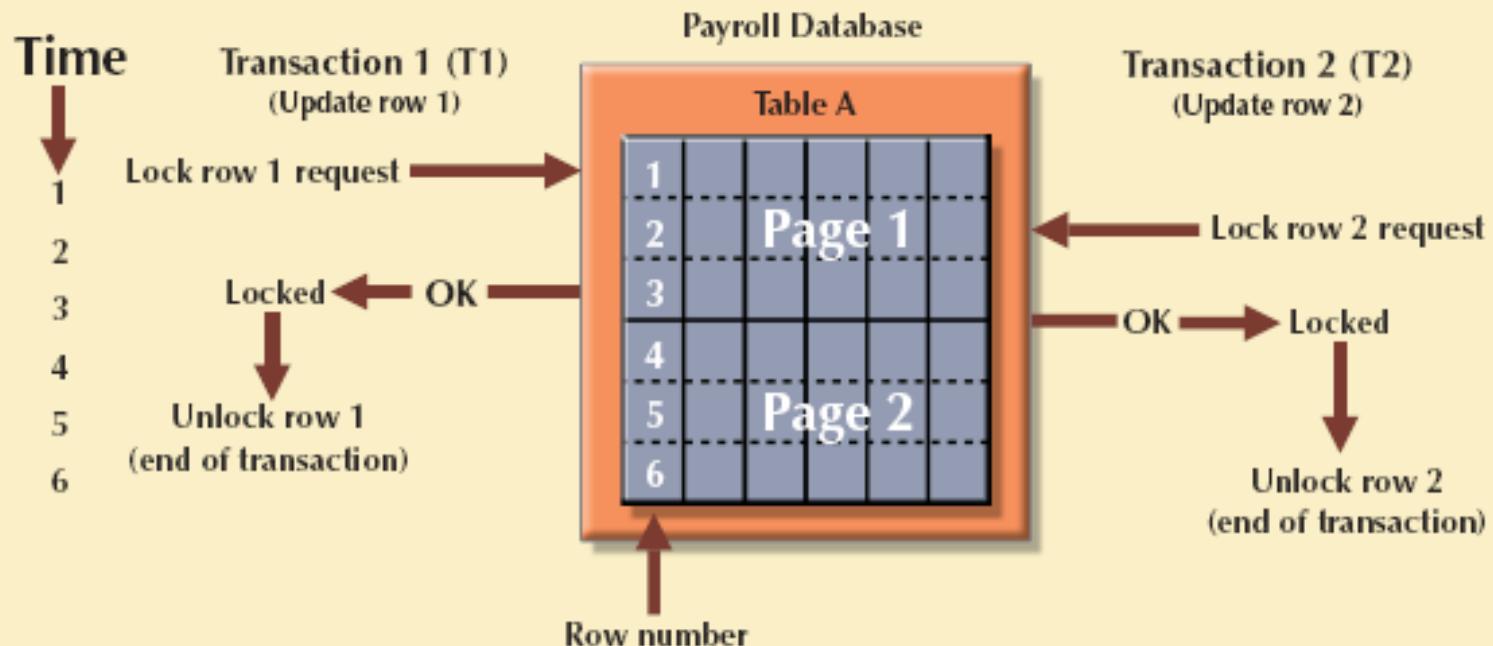
SOURCE: Course Technology/Cengage Learning



Edit with WPS Office

# ROW-LEVEL LOCK

FIGURE 13.6 An example of a row-level lock



SOURCE: Course Technology/Cengage Learning



Edit with WPS Office

# BINARY LOCKS

## ☒ Binary lock

- ☒ Two states: locked (1) or unlocked (0)

☒ TABLE  
13.12

An Example of a Binary Lock

TIME	TRANSACTION	STEP	STORED VALUE
1	T1	Lock PRODUCT	
2	T1	Read PROD_QOH	15
3	T1	$PROD\_QOH = 15 + 10$	
4	T1	Write PROD_QOH	25
5	T1	Unlock PRODUCT	
6	T2	Lock PRODUCT	
7	T2	Read PROD_QOH	23
8	T2	$PROD\_QOH = 23 - 10$	
9	T2	Write PROD_QOH	13
10	T2	Unlock PRODUCT	



Edit with WPS Office

# SHARED/EXCLUSIVE LOCKS

## ⊕ Exclusive lock

- ⊗ Access is specifically reserved for transaction that locked object
- ⊗ Must be used when potential for conflict exists (when a transaction wants to write data)

## ⊕ Shared lock

- ⊗ Issued when a transaction wants to read data and no exclusive lock is held on that data item
- ⊗ Concurrent transactions are granted read access on basis of a common lock
- ⊗ No conflict as long as all the concurrent transactions are read-only



Edit with WPS Office

# TWO-PHASE LOCKING

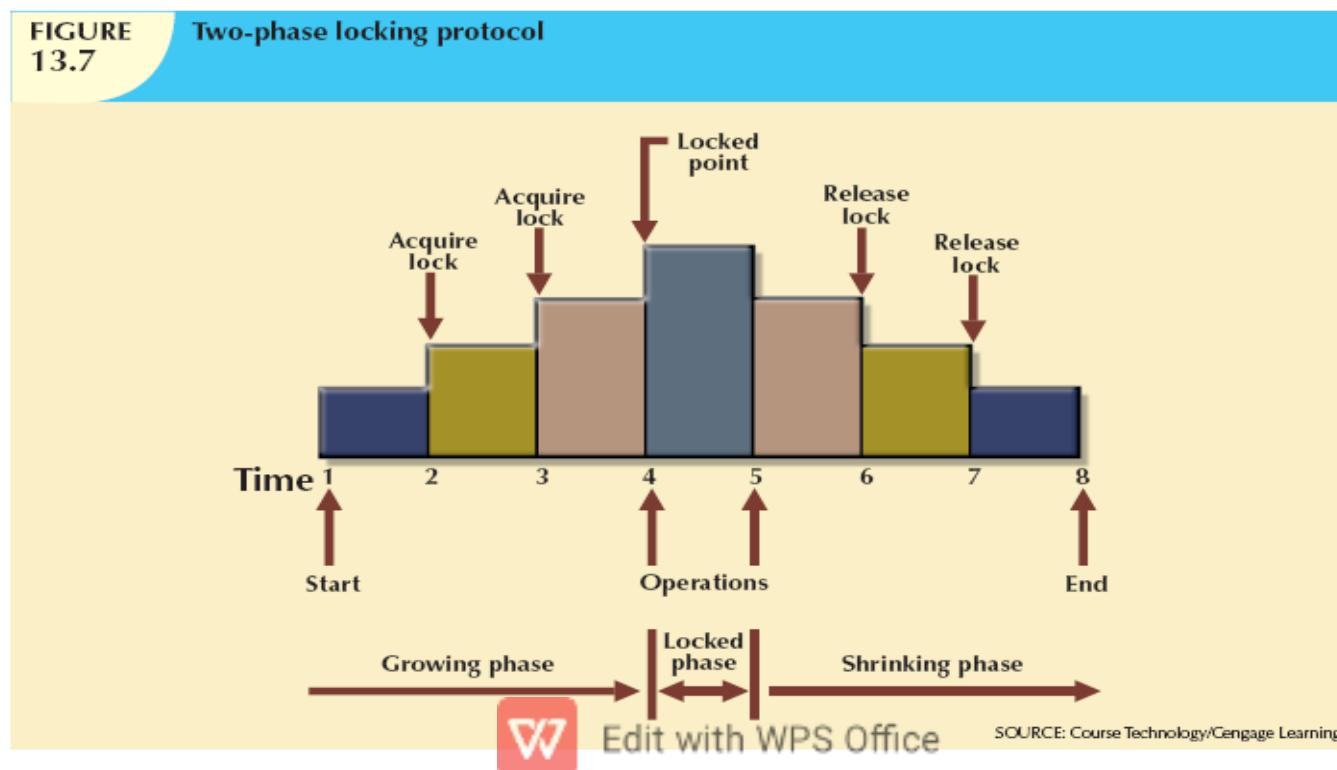
- ☒ Defines how transactions acquire and relinquish locks
- ☒ Guarantees serializability, but does not prevent deadlocks
- ☒ Two phases:
  - ☒ Growing phase
    - ☐ Transaction acquires all required locks without unlocking any data
  - ☒ Shrinking phase
    - ☐ Transaction releases all locks and cannot obtain any new lock



Edit with WPS Office

# TWO-PHASE LOCKING (CONT’)

- ☒ Governed by the following rules:
  - ☒ Two transactions cannot have conflicting locks
  - ☒ No unlock operation can precede a lock operation in the same transaction
  - ☒ No data are affected until all locks are obtained



# DEADLOCKS

- ☒ Condition that occurs when two transactions wait for each other to unlock data
- ☒ Possible only if one of the transactions wants to obtain an exclusive lock on a data item
  - ☒ No deadlock condition can exist among shared locks
- ☒ Example:
  - ☒ T1 = access data items X then Y
  - ☒ T2 = access data items Y then X
  - ☒ T1 acquires the lock on X first and waits for T2 to release the lock on Y to commit the transaction
  - ☒ T2 acquires the lock on Y first and waits for T1 to release the lock on X to commit the transaction



Edit with WPS Office

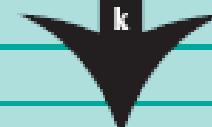
# DEADLOCKS - EXAMPLE

TABLE  
13.13

How a Deadlock Condition Is Created

TIME	TRANSACTION	REPLY	LOCK STATUS	
0			Data X	Data Y
1	T1:LOCK(X)	OK	Unlocked	Unlocked
2	T2: LOCK(Y)	OK	Locked	Unlocked
3	T1:LOCK(Y)	WAIT	Locked	Locked
4	T2:LOCK(X)	WAIT	Locked	Locked
5	T1:LOCK(Y)	WAIT	Locked	Locked
6	T2:LOCK(X)	WAIT	Locked	Locked
7	T1:LOCK(Y)	WAIT	Locked	Locked
8	T2:LOCK(X)	WAIT	Locked	Locked
9	T1:LOCK(Y)	WAIT	Locked	Locked
...	.....	.....	.....	.....
...	.....	.....	.....	.....
...	.....	.....	.....	.....
...	.....	.....	.....	.....

Deadlock



Edit with WPS Office

# DEADLOCKS (CONT’)

- ☒ Three techniques to control deadlock:
  - ☒ Prevention
  - ☒ Detection
  - ☒ Avoidance
- ☒ Choice of deadlock control method depends on database environment
  - ☒ Low probability of deadlock; detection recommended
  - ☒ High probability; prevention recommended



Edit with WPS Office

# CONCURRENCY CONTROL WITH TIME STAMPING METHODS

- ☒ Assigns global unique time stamp to each transaction
- ☒ Produces explicit order in which transactions are submitted to DBMS
- ☒ Uniqueness
  - ☒ Ensures that no equal time stamp values can exist
- ☒ Monotonicity
  - ☒ Ensures that time stamp values always increase



Edit with WPS Office

# WAIT/DIE AND WOUND/WAIT SCHEMES

## Wait/die

- Older transaction waits and younger is rolled back and rescheduled

## Wound/wait

- Older transaction rolls back younger transaction and reschedules it

TABLE  
13.14

Wait/Die and Wound/Wait Concurrency Control Schemes

TRANSACTION REQUESTING LOCK	TRANSACTION OWNING LOCK	WAIT/DIE SCHEME	WOUND/WAIT SCHEME
T1 (11548789)	T2 (19562545)	<ul style="list-style-type: none"><li>T1 waits until T2 is completed and T2 releases its locks.</li></ul>	<ul style="list-style-type: none"><li>T1 preempts (rolls back) T2.</li><li>T2 is rescheduled using the same timestamp.</li></ul>
T2 (19562545)	T1 (11548789)	<ul style="list-style-type: none"><li>T2 dies (rolls back).</li><li>T2 is rescheduled using the same timestamp.</li></ul>	<ul style="list-style-type: none"><li>T2 waits until T1 is completed and T1 releases its locks.</li></ul>



Edit with WPS Office

# **DATABASE RECOVERY MANAGEMENT**

- ☒ Restores database to previous consistent state
- ☒ Based on atomic transaction property
  - ☒ All portions of transaction are treated as single logical unit of work
  - ☒ All operations are applied and completed to produce consistent database
- ☒ If transaction operation cannot be completed:
  - ☒ Transaction aborted
  - ☒ Changes to database are rolled back



Edit with WPS Office

# TRANSACTION RECOVERY

- ☒ Write-ahead-log protocol
  - ☒ Ensures transaction logs are written before data is updated
- ☒ Redundant transaction logs
  - ☒ Ensure physical disk failure will not impair ability to recover
- ☒ Buffers
  - ☒ Temporary storage areas in primary memory
- ☒ Checkpoints
  - ☒ Operations in which DBMS writes all its updated buffers to disk



Edit with WPS Office

# TRANSACTION RECOVERY (CONT’)

- ☒ Deferred-write technique
  - ☒ Only transaction log is updated
- ☒ Recovery process: identify last checkpoint
  - ☒ If transaction committed before checkpoint:
    - ☐ Do nothing
  - ☒ If transaction committed after checkpoint:
    - ☐ Use transaction log to redo the transaction
  - ☒ If transaction had ROLLBACK operation:
    - ☐ Do nothing



Edit with WPS Office

# TRANSACTION RECOVERY (CONT)

- ☒ Write-through technique
  - ☒ Database is immediately updated by transaction operations during transaction's execution
- ☒ Recovery process: identify last checkpoint
  - ☒ If transaction committed before checkpoint:
    - ☐ Do nothing
  - ☒ If transaction committed after last checkpoint:
    - ☐ DBMS redoes the transaction using “after” values
  - ☒ If transaction had ROLLBACK or was left active:
    - ☐ DBMS undoes the operation using the “before” values in the transaction log



Edit with WPS Office

# TRANSACTION RECOVERY - EXAMPLE

TABLE  
13.15

A Transaction Log for Transaction Recovery Examples

TRL ID	TRX NUM	PREV PTR	NEXT PTR	OPERATION	TABLE	ROW ID	ATTRIBUTE	BEFORE VALUE	AFTER VALUE
341	101	Null	352	START	****Start Transaction				
352	101	341	363	UPDATE	PRODUCT	54778-2T	PROD_QOH	45	43
363	101	352	365	UPDATE	CUSTOMER	10011	CUST_BALANCE	615.73	675.62
365	101	363	Null	COMMIT	**** End of Transaction				
397	106	Null	405	START	****Start Transaction				
405	106	397	415	INSERT	INVOICE	1009			1009,10016, ...
415	106	405	419	INSERT	LINE	1009,1			1009,1, 89-WRE-Q,1, ...
419	106	415	427	UPDATE	PRODUCT	89-WRE-Q	PROD_QOH	12	11
423				CHECKPOINT					
427	106	419	431	UPDATE	CUSTOMER	10016	CUST_BALANCE	0.00	277.55
431	106	427	457	INSERT	ACCT_TRANSACTION	10007			1007,18-JAN-2012, ...
457	106	431	Null	COMMIT	**** End of Transaction				
521	155	Null	525	START	****Start Transaction				
525	155	521	528	UPDATE	PRODUCT	2232/QWE	PROD_QOH	6	26
528	155	525	Null	COMMIT	**** End of Transaction				
***** C *R*A* S* H *****									



Edit with WPS Office

# SUMMARY

- ☒ Transaction: sequence of database operations that access database
  - ☒ Logical unit of work
    - ☐ No portion of transaction can exist by itself
  - ☒ Five main properties: atomicity, consistency, isolation, durability, and serializability
- ☒ COMMIT saves changes to disk
- ☒ ROLLBACK restores previous database state
- ☒ SQL transactions are formed by several SQL statements or database requests



Edit with WPS Office

# SUMMARY (CONT'D.)

- ☒ Transaction log keeps track of all transactions that modify database
- ☒ Concurrency control coordinates simultaneous execution of transactions
- ☒ Scheduler establishes order in which concurrent transaction operations are executed
- ☒ Lock guarantees unique access to a data item by transaction
- ☒ Two types of locks: binary locks and shared/exclusive locks



Edit with WPS Office

# SUMMARY (CONT'D.)

- ☒ Serializability of schedules is guaranteed through the use of two-phase locking
- ☒ Deadlock: when two or more transactions wait indefinitely for each other to release lock
- ☒ Three deadlock control techniques: prevention, detection, and avoidance
- ☒ Time stamping methods assign unique time stamp to each transaction
  - ☒ Schedules execution of conflicting transactions in time stamp order
- ☒ Database recovery restores database from given state to previous consistent state



Edit with WPS Office